

Toward Hardware-Adaptive Pareto Fronts for CPU-Only LLM Inference: An Empirical Study of llama.cpp Across Four Heterogeneous Devices

Yassine

Master's Thesis Report

August 18, 2026

Abstract

Running large language models without a GPU is, for most of the world’s compute, the only option available: aging desktops, laptops bought for office work, single-board computers, and embedded devices vastly outnumber machines with a discrete accelerator. This raises a practical question that is easy to state and hard to answer in general: for a *given* CPU, what combination of model, quantization format, thread count, batch size, and KV-cache precision actually maximizes throughput, and at what cost in output quality? We argue that `llama.cpp` is the right vehicle for studying this question empirically, not because it is the fastest possible CPU inference engine in any single configuration, but because it is the most *parametrically exposed* one: a single binary, `llama-bench`, sweeps threads, context depth, batch size, KV-cache quantization, Flash Attention, mmap/NUMA behavior, and memory-mapping policy without requiring the user to touch source code, and it runs unmodified on architectures as different as an AMD Zen 3 laptop chip, a decade-old Intel Pentium Gold desktop part, and a Cortex-A57 Jetson Nano. We built a benchmarking harness around `llama-bench` and ran it on four machines (three of which produced usable data), spanning roughly a $70\times$ range in single-thread generation throughput, recovering three concrete, reproducible effects: (1) prompt processing is compute-bound and follows Gunther’s Universal Scalability Law, saturating near each CPU’s physical (not logical) core count; (2) single-batch token generation is memory-bandwidth-bound and can show *retrograde* scaling, i.e. adding threads makes it slower, most dramatically on bandwidth-starved processors; and (3) quantization format matters independently of bit-width – a 3-bit K-quant was measured to be *slower* than a 4-bit K-quant on the same hardware, and quantizing the KV cache was measured to be a net loss at all tested context depths on a compute-rich laptop. We further document a fourth effect: **Flash Attention is a performance regression on CPU at long context**, with generation throughput dropping by a factor of three at 15K-token context when Flash Attention is enabled on our test hardware. From these effects we derive, rather than fit, a roofline-style throughput model with a small number of per-device calibration constants, and we show how it extends naturally to Mixture-of-Experts (MoE) models via a “bytes touched per token” argument that predicts MoE should be disproportionately advantaged on bandwidth-starved CPUs – a prediction we lay out as the clearest next experiment, since no MoE model was benchmarked to completion on our own hardware in this round. We also built an interactive web calculator (https://yassine-hattay.github.io/LLMs_inference_on_CPU/index.html) implementing this model so practitioners can explore predicted throughput for their own hardware traits; it is a client-side prediction tool, not a live benchmark, and its predictions beyond the three directly-calibrated devices are extrapolations, not measurements. We release the benchmarking harness, the fitted constants, and the calculator so that additional devices can be folded in without re-deriving the model.

Contents

Introduction	1
Motivation: the GPU-less majority	1
Why llama.cpp	2
Dense vs. Mixture-of-Experts models	2
Contributions	3
1 Related Work	4
Introduction	4
1.1 CPU LLM Inference Frameworks	4
1.2 Roofline Modeling and Scalability Laws	5
1.3 Quantization for LLMs	5
1.4 Mixture-of-Experts Architectures	6
1.5 Edge and Embedded LLM Deployment	6
1.6 Positioning of This Work	7
Conclusion	9
2 Instrumentation and Data Collection	10
Introduction	10
2.1 Hardware fleet	10
2.2 Hardware micro-architecture details	11
2.3 Models tested	11
2.4 The measurement instrument	12
2.5 Experimental design: what was swept, and why	12
2.6 Resumability: design and an instructive failure	13
2.7 Measurement protocol	14
Conclusion	16
3 Results	17
Introduction	17
3.1 Thread scaling follows the Universal Scalability Law, not simple Amdahl's law	17
3.2 Model size determines which regime you're in	19
3.3 Flash Attention is a performance regression on CPU at long context	21
3.4 KV-cache quantization is not a free lunch	22
3.5 Quantization format is not just bit-width	23
3.6 Reference: quantization quality trade-offs from the literature	25
3.7 Raspberry Pi: attempted but not productive	25
3.8 The MoE bandwidth argument remains untested on our hardware	26

Conclusion	26
4 A Pareto-Front Model and Practical Recommendations	28
Introduction	28
4.1 What we can and cannot claim	28
4.2 A roofline model with USL thread scaling	29
4.3 KV-cache as a second bandwidth term	30
4.4 Worked example: calibrating the model for the Ryzen 5600H	30
4.5 Defining the Pareto front	31
4.6 Practical Deployment Recommendations	31
Conclusion	32
5 Discussion, Limitations, and Future Work	33
Introduction	33
5.1 Discussion	33
5.2 Future Work	35
5.2.1 Near-Term: Completing the Current Study	36
5.2.2 Longer-Term Directions	36
Conclusion	38
6 Conclusion	40

List of Figures

3.1	Ryzen 5600H: prefill saturates near the physical core count; generation retrogrades past 4 threads.	19
3.2	Pentium G5420: prefill saturates by 4 threads; generation retrogrades monotonically.	19
3.3	Jetson Nano: prefill scales for both models; generation only scales for the smaller one.	20
3.4	Ryzen 5600H, 1 thread: every quantized-KV configuration was slower than f16, worsening with context depth.	23
3.5	Pentium G5420, 1 thread: quantization <i>format</i> , not just bit-width, determines CPU throughput.	24

List of Tables

1.1	Related work compared against the axes this thesis covers. “Formal model?” means the work proposes an equation or fitted curve, not just raw measurements. No prior work we found combines a roofline-style throughput model, USL thread scaling, quantization-format effects, and an MoE bandwidth argument in one framework – but as the table also makes clear, our own MoE column is a hypothesis, not a result, which is the most direct gap between this thesis’s ambition and what it has actually shown. . . .	8
2.1	Benchmarked hardware. All runs used <code>-ngl 0</code> (no GPU offload); the laptop’s integrated Radeon graphics and the Jetson’s Maxwell GPU were both left idle so every number in this paper is pure CPU throughput. Memory bandwidth figures are theoretical peak; effective sustained bandwidth (Section 4.1) is typically 60–80% of this. The Raspberry Pi 3B was included in the hardware fleet but its benchmark run did not complete – the output file it produced contains zero rows – so it does not appear in any results below; see Section 5.1. . . .	11
2.2	Micro-architecture details (manufacturer specifications). Note the Pentium’s lack of AVX-512/VNNI and the Jetson’s cacheless L3 – both are relevant to interpreting the quantization and context-depth results. . . .	11
2.3	Models staged for benchmarking. MoE models’ “active” parameter counts are per-token; their total RAM footprint is determined by the full parameter count. Only the dense models appear in our actual results (Section 3); the MoE models were staged but not benchmarked to completion. . . .	12
3.1	USL fits. Note that $\kappa_g > \kappa_p$ in all cases, confirming that generation is more contention-sensitive than prefill. The Pentium’s generation $N^* = 1$ means single-threaded is optimal. . . .	20
3.2	Ryzen 5600H raw thread-scaling data, all context depths (<code>avg_ts</code> , tok/s). . . .	20
3.3	Jetson Nano raw model-size comparison data (<code>avg_ts</code> , tok/s). . . .	21
3.4	Flash Attention vs. standard attention on Ryzen 5600H at 1 thread. At short context (depth=0), Flash Attention is neutral. At long context (depth=15872), Flash Attention is a $3\times$ <i>slowdown</i> for generation. . . .	21
3.5	KV-cache quantization on Ryzen 5600H at 1 thread, across context depths. KV quantization is a net loss at <i>all</i> tested context depths on this hardware, and the loss grows with context length. . . .	22
3.6	Quantization format comparison on Pentium G5420. Q4_0 is $3.6\times$ faster than Q4_K_M at 1 thread despite the same bit-width. Q3_K_M is smaller than Q4_K_M but slower. . . .	24

3.7	The only externally-sourced MoE data point we could independently verify; roughly $80\times$ larger by active-parameter count than the MoE models staged for this study.	26
5.1	Data coverage by machine and phase; every unattributed claim in this thesis is drawn only from rows marked complete.	34

Introduction

Motivation: the GPU-less majority

The dominant narrative around large language model (LLM) inference assumes a GPU. Benchmark leaderboards report tokens-per-second on H100s and consumer RTX cards; deployment guides assume CUDA is present. But a very large share of the machines people actually own – laptops purchased for browsing and office work, hand-me-down desktops, Raspberry Pis, and low-power edge boards like the NVIDIA Jetson Nano – have no usable GPU for inference at all, or have one too memory-constrained to hold even a small model. For that hardware, the only path to running an LLM locally is CPU inference, and the only question that matters practically is: *given this specific chip, this specific amount of RAM, and this specific model, which combination of settings gets me the most tokens per second at output quality I can accept?*

This thesis does not propose a new inference engine or a new quantization scheme. It instead treats an existing, widely deployed one – `llama.cpp` – as a measurement instrument, and asks how far a small, heterogeneous fleet of consumer and embedded machines can get us toward answering that question in a principled, reusable way.

The practical motivation for this study stems from the configuration complexity faced by practitioners operating on constrained hardware. When attempting to deploy a local language model on a standard consumer laptop or single-board computer lacking a discrete GPU, users are confronted with a high-dimensional, non-intuitive parameter space. They must navigate an extensive array of execution flags (e.g., `-t`, `-n`, `-b`, `-ub`, `-c`, `-ctk`, `-ctv`, `-fa`, `-mmp`) and select from a diverse ecosystem of quantized model variants (e.g., `Q4_K_M`, `Q4_0`, `Q3_K_S`, `Q8_0`, `IQ4_NL`). Without a principled, hardware-aware framework to guide these decisions, practitioners are often forced to rely on trial-and-error or anecdotal community advice, frequently resulting in suboptimal performance or execution failures. Recent work has begun to characterize this space from several angles.

RooflineBench [1] applies roofline analysis to on-device LLMs but focuses primarily on mobile SoCs with unified memory. SMEPilot [2] evaluates roofline models for CPU LLM execution on ARM platforms equipped with Scalable Matrix Extensions, while HeteroMosaic [3] formulates a heterogeneous roofline model across integrated GPUs and NPUs, treating the CPU primarily as a fallback. The quantization dimension has been studied in isolation: a unified evaluation of `llama.cpp` quantization formats on Llama-3.1-8B [4] characterizes perplexity-speed trade-offs for a single architecture. Community benchmarking efforts on forums such as `r/LocalLLaMA` [5] provide extensive raw through-

put numbers across hardware configurations but do not synthesize them into a predictive model. What remains missing is a unified framework that jointly models thread scaling, memory bandwidth saturation, quantization format overhead, and KV-cache effects across the full spectrum of consumer and embedded CPUs – from compute-rich laptop processors to bandwidth-starved single-board computers – and that extends to a principled prediction for Mixture-of-Experts architectures. This work is a step toward that framework.

Why `llama.cpp`

Three properties made `llama.cpp` the natural choice for this study over alternatives such as vLLM, MLC-LLM, ONNX Runtime, or PyTorch with `bitsandbytes`:

1. **Breadth of CPU kernel support.** `llama.cpp`'s GGML backend ships hand-written SIMD kernels for x86 (SSE/AVX/AVX2/AVX-512/VNNI), ARM NEON, and generic scalar fallbacks, and it selects among them at runtime. This is precisely what let the same benchmark script, unmodified except for a thread-count override, run on a Zen 3 laptop, a Coffee-Lake-era Pentium, and a Cortex-A57 SBC.
2. **A dense, orthogonal parameter surface exposed through one tool.** `llama-bench` accepts thread count (`-t`), context depth (`-d`), prompt/generation length (`-p/-n`), batch and micro-batch size (`-b/-ub`), KV-cache quantization for keys and values independently (`-ctk/-ctv`), Flash Attention on/off (`-fa`), memory mapping (`-mmp`), and number of GPU-offloaded layers (`-ngl`, held at 0 throughout this study) as simple flags, and emits structured JSONL. No other CPU inference stack exposes this many independent axes through a single stable command-line interface.
3. **The GGUF quantization ecosystem.** Pre-quantized weights for essentially every open model of interest (dense and MoE) are available in GGUF, at bit-widths from 2 to 8, including both the original “legacy” formats (Q4_0, Q5_0, ...) and the newer super-block “K-quants” (Q3_K_M, Q4_K_M, ...), which let us treat quantization format itself as an experimental variable rather than a fixed implementation detail.

Dense vs. Mixture-of-Experts models

The models staged for this study span both major architecture families in current open-weight LLMs:

- **Dense models** (Qwen2.5-0.5B/1.5B/7B-Instruct, Qwen3-8B, Meta-Llama-3.1-8B-Instruct, Phi-4) activate *every* parameter on *every* forward pass. Their compute and memory-traffic cost per generated token is therefore fixed by their total parameter count and quantization bit-width, with no runtime variability.

- **Mixture-of-Experts (MoE) models** (Qwen1.5-MoE-A2.7B-Chat, DeepSeek-V2-Lite, DeepSeek-Coder-V2-Lite-Base) replace the feed-forward block with a bank of experts and a lightweight router that activates only a handful of them per token. A model can therefore have, e.g., 14B *total* parameters but only ~ 2.7 B *active* parameters per token – the rest of the weights exist on disk and in RAM, but are not read on every forward pass.

For GPU inference, where the whole model is typically resident in fast VRAM and compute (not bandwidth) is often the binding constraint at low batch sizes, this distinction matters less. For *CPU* inference at batch size 1 – the realistic single-user case this study targets – generation is, as Section 3 shows directly, memory-bandwidth-bound. Under that regime, a dense model must stream its entire quantized weight set from RAM for every single token it emits, while an MoE model only has to stream the weights of the experts its router actually selected. This is the theoretical reason MoE is the focal point of this line of work (Section 4.1): it acts as a lever for decoupling generation speed from total model capacity in exactly the regime where CPU inference lives. By analyzing the memory-bandwidth constraints observed in our dense-model benchmarks, we can directly derive and formalize the performance advantage MoE architectures provide on these constrained systems.

Contributions

1. An empirical characterization (Section 3) of thread scaling, KV-cache quantization overhead, Flash Attention behavior, and quantization-format effects across four real, heterogeneous CPUs (three of which produced usable data), fit where the data supports it to Gunther’s Universal Scalability Law rather than simple Amdahl’s law, because the data shows genuine retrograde (not just saturating) scaling.
2. A roofline-style throughput model (Section 4.1) with per-device calibration constants that formalizes *why* the observed effects occur, extends to a bytes-touched-per-token argument for MoE models, and defines a Pareto front over (speed, quality, memory) that can be populated directly by local benchmarks or community-sourced data.
3. Practical deployment recommendations (Section 4.6) distilled from the data, giving concrete guidance for users running `llama.cpp` on consumer hardware.
4. A discussion of the model’s scope and boundaries (Section 5.1), outlining how the per-device calibration framework can be extended to broader hardware fleets and emerging MoE architectures.

Chapter 1

Related Work

Introduction

This thesis draws on four bodies of prior work: CPU-oriented LLM inference frameworks, roofline and scalability modeling, post-training quantization, and Mixture-of-Experts architectures. None of these individually addresses the question this thesis asks – what determines optimal `llama.cpp` configuration on a specific, resource-constrained CPU – but each supplies a piece of the machinery this thesis assembles. This chapter reviews each body of work in turn, adds a fifth section on edge and embedded LLM deployment specifically (the hardware class most of this thesis’s own fleet belongs to), and closes by positioning this thesis explicitly against the six most directly comparable efforts in Table 1.1.

1.1 CPU LLM Inference Frameworks

The landscape of LLM inference engines is dominated by GPU-first frameworks. **vLLM** [6] introduced PagedAttention and is the de facto standard for high-throughput GPU serving, but its CPU backend is a secondary concern and does not expose the same parametric sweep surface as `llama.cpp`. **MLC-LLM** [7] targets cross-platform deployment including WebGPU and mobile, but its CPU path relies on TVM-compiled kernels that are less transparent to end users. **ONNX Runtime** with the **OpenVINO** or **ONNX Runtime GenAI** extensions provides a well-optimized CPU path, but its quantization story is oriented around INT4/INT8 symmetric formats rather than the GGUF ecosystem’s asymmetric K-quants. **PyTorch** with **bitsandbytes** or **auto-gptq** is the research community’s default, but its CPU performance is typically an order of magnitude slower than `llama.cpp` on the same hardware due to the overhead of the PyTorch runtime and lack of hand-tuned CPU kernels.

`llama.cpp` occupies a unique niche: it is the only framework that (a) runs unmodified on x86, ARM, and RISC-V, (b) exposes a dense parameter surface through a single CLI, and (c) has a thriving ecosystem of pre-quantized GGUF models on Hugging Face. This

makes it the natural choice for empirical studies of CPU inference, and it is the framework used in the growing body of community benchmarking literature [8, 5]. Crucially for this thesis’s methodology (Chapter 2), **llama-bench**’s structured JSONL output and its independent flags for thread count, KV-cache quantization, and Flash Attention are what make a controlled, single-variable-at-a-time sweep practical at all; none of the alternative frameworks above expose that many orthogonal axes through one stable interface, which is as much a methodological finding of this review as it is a justification for the tool choice.

1.2 Roofline Modeling and Scalability Laws

The **roofline model** [9] is a foundational tool in high-performance computing, visualizing an application’s performance relative to hardware peaks in compute (FLOP/s) and memory bandwidth (bytes/s). It has been applied to GPUs, TPUs, and specialized accelerators, but its application to CPU-based LLM inference – where the “compute peak” is a moving target depending on SIMD width and the “bandwidth peak” is the sustained (not burst) DRAM throughput – is relatively under-explored. Recent work such as **RooflineBench** [1] has begun to apply roofline analysis to on-device LLMs, but focuses primarily on mobile SoCs with unified memory, where CPU and GPU compete for the same memory pool in a way that does not generalize cleanly to a discrete desktop or laptop CPU with its own dedicated DRAM controller. **SMEPilot** [2] and **HeteroMosaic** [3], discussed further in Section 1.5, extend the roofline lens to newer ARM vector extensions and to heterogeneous CPU/iGPU/NPU systems respectively, but both treat the CPU as either a narrow special case (a single modern ISA extension) or a fallback path behind a faster accelerator, rather than as the sole target of the analysis.

Gunther’s Universal Scalability Law (USL) [10] extends classical Amdahl’s law by adding a coherency term $\kappa N(N - 1)$ that captures crosstalk and contention penalties, on top of Amdahl’s own serialization term σ . USL has been applied extensively outside of machine learning – most notably to database systems, web servers, and distributed transaction processing, domains where adding worker threads or nodes eventually causes throughput to *decline*, not just plateau, because of lock contention or cache-coherency traffic. Its application to LLM inference thread scaling is, to our knowledge, novel: the closest existing work (RooflineBench, SMEPilot) models the compute/bandwidth boundary but does not model contention between threads explicitly, so it predicts saturation but not the downturn itself. Classical Amdahl’s law predicts throughput that rises and asymptotically flattens; USL predicts a peak and possible downturn – exactly what we observe in our generation data on bandwidth-starved CPUs (Section 3.1), and the reason this thesis fits USL rather than Amdahl’s law wherever the data supports it.

1.3 Quantization for LLMs

Post-training quantization (PTQ) for LLMs has seen rapid development. **GPTQ** [11] and **AWQ** [12] produce INT4/INT3 symmetric quantizations optimized for GPU kernels, us-

ing layer-wise reconstruction error minimization and activation-aware scaling respectively; both are designed around GPU tensor-core throughput and are not the formats this thesis measures, but they represent the GPU-side counterpart to the CPU-side GGUF ecosystem and are useful context for why the two ecosystems have diverged. The GGUF ecosystem’s **K-quants** [13] take a different approach: super-block bit-packing with per-block scales, optimized for perplexity-per-bit rather than kernel throughput. This distinction is critical for CPU inference: K-quants achieve better quality at a given bit-width, but their unpack kernels are more complex and can be slower on CPUs without AVX-512/VNNI. A separate, single-architecture survey of this trade-off, “**Which Quantization...**” [4], characterizes the perplexity-versus-speed frontier for Llama-3.1-8B specifically; it establishes that the trade-off exists but does not test whether it holds across different, more constrained CPU micro-architectures. Our data in Section 3.5 directly demonstrates this trade-off on hardware that survey does not cover, and shows that the direction of the trade-off can flip entirely on legacy CPUs lacking modern vector instructions – a finding absent from the existing quantization literature.

1.4 Mixture-of-Experts Architectures

MoE models have emerged as a way to scale model capacity without proportionally scaling compute. **Mixtral 8x7B** [14], **DeepSeek-V2** [15], and **Qwen1.5-MoE** [16] all use sparse routing to activate only a fraction of parameters per token. On GPU, MoE models are primarily a capacity play (more parameters, same VRAM), because the full parameter set is typically resident in VRAM regardless of how many experts are active on a given forward pass, so the benefit is measured in what a model can *know*, not how fast it responds. On CPU, as we argue in Section 4.1, MoE models are instead a *bandwidth* play: because weights must be streamed from RAM on every token rather than sitting in fast on-chip memory, activating fewer experts directly reduces the bytes moved per token, which is exactly the bottleneck this thesis identifies as binding for CPU generation (Section 3.2). This insight – that MoE’s value proposition flips from a capacity argument on GPU to a bandwidth argument on CPU – is a central contribution of this work, and it is, as far as we have been able to establish, not stated explicitly in the systems literature on either MoE architectures or CPU inference individually; it emerges only from combining the two.

1.5 Edge and Embedded LLM Deployment

A distinct and rapidly growing body of practitioner-facing work targets LLM inference specifically on edge and single-board hardware – the class of device that three of this thesis’s four target machines (the Jetson Nano, the Pentium G5420, and the attempted Raspberry Pi 3B) belong to. Jeff Geerling’s publicly maintained AI benchmarking repository [17] and a 2026 survey of Raspberry Pi 5 LLM performance [18] both report empirical throughput numbers for 7–8B-class quantized models on Raspberry Pi hardware, and are the closest published analogue to the Raspberry Pi data this thesis was unable to collect itself (Section 3.7). Neither, however, fits a predictive model to its measurements the

way this thesis’s roofline-USL framework does (Section 4.1); they report throughput as a function of configuration but do not explain *why* a given configuration is fast or slow in terms of the compute/bandwidth boundary. **HeteroMosaic** [3] is closer in spirit – it formalizes a heterogeneous roofline model across integrated GPUs, NPUs, and CPUs on edge SoCs – but it treats the CPU path as a fallback to be modeled only well enough to know when to route work elsewhere, rather than as the object of study in its own right. This thesis’s contribution relative to this whole literature is narrow but specific: a CPU-only roofline-USL model, calibrated per device, applied to exactly the low-end and embedded hardware this literature otherwise only benchmarks empirically.

1.6 Positioning of This Work

Our work sits at the intersection of these threads: we apply roofline modeling and USL to empirical `llama.cpp` data across heterogeneous consumer CPUs, we characterize the quantization-format and KV-cache trade-offs that prior work has not systematically measured, and we derive a predictive argument for MoE advantage on bandwidth-starved hardware. To our knowledge, no prior work has combined all four of these elements into a single coherent framework. Table 1.1 makes this comparison explicit, scoring each related effort against the axes this thesis covers so the overlap – and, just as importantly, the gaps neither we nor prior work close – are visible at a glance.

Work	Focus	Hardware scope	Formal model?	MoE?	Relation to this thesis
RooflineBench [1]	Roofline analysis for on-device LLMs	Mobile SoCs, unified memory	Yes	No	Same roofline lens; we target discrete consumer/embedded CPUs instead of mobile unified memory.
SMEPilot [2]	Roofline modeling for ARM CPUs with Scalable Matrix Extensions	Modern ARM SME	Yes	No	Shares the roofline+ARM focus; targets a newer ISA extension absent from our fleet (Cortex-A57/A53).
HeteroMosaic [3]	Heterogeneous roofline across NPU/iGPU/CPU	Edge SoCs with NPU+iGPU	Yes	No	CPU is treated as a fallback path there; it is our primary and only target here.
“Which Quantization...” [4]	Perplexity-vs-speed survey of llama.cpp quant formats	Single architecture (Llama-3.1-8B)	No (empirical survey)	No	We extend the quantization-format question across multiple, more constrained CPUs and tie it into a throughput model (Section 3.5).
KTransformers docs [19]	CPU/GPU heterogeneous serving for large MoE models	Server-class socket Xeon	No (engineering benchmark)	Yes	Our one verifiable external MoE data point (Section 3.8); different model scale and server-class hardware, not consumer/embedded.
Community forums (r/LocalLLaMA, etc.) [5]	Crowd-sourced throughput reports	Broad, unsystematic	No	Sporadic	Rich raw numbers but no synthesis into a predictive model; a partial substitute for measurements we could not collect ourselves (Section 3.7).
This thesis	Roofline + USL predictive model, calibrated per device, for llama.cpp on consumer/embedded CPUs	Laptop, low-end desktop, embedded ARM SBC (3 devices with usable data)	Yes (roofline + USL)	Hypothesis, explicitly untested locally	

Table 1.1: Related work compared against the axes this thesis covers. “Formal model?” means the work proposes an equation or fitted curve, not just raw measurements. No prior work we found combines a roofline-style throughput model, USL thread scaling, quantization-format effects, and an MoE bandwidth argument in one framework – but as the table also makes clear, our own MoE column is a hypothesis, not a result, which is the most direct gap between this thesis’s ambition and what it has actually shown.

Conclusion

The gap this chapter identifies is narrow but real: existing roofline-style analyses of on-device LLM inference target mobile SoCs or newer ARM extensions, existing quantization surveys target a single architecture, existing MoE literature treats CPU behavior only in passing or on server-class hardware, and existing edge-deployment benchmarks report throughput without a predictive model behind it. No source we reviewed combines a fitted scalability law, a roofline throughput model, a systematic quantization-format comparison, and an explicit CPU-side reframing of the MoE value proposition into one account of consumer and embedded CPU inference. Table 1.1 makes this gap, and the corresponding limits of this thesis’s own contribution, explicit rather than implicit. The next chapter describes the instrumentation built to collect the data that fills it.

Chapter 2

Instrumentation and Data Collection

Introduction

This chapter describes how the data behind Chapter 3 was actually collected: the hardware fleet and its micro-architectural details, the models staged for benchmarking, the design logic behind what was swept and why, the resumable execution scheme that let unattended multi-hour runs survive interruption on non-dedicated hardware, and the measurement protocol used to turn raw `llama-bench` output into the throughput numbers reported later. It is organized to be read either end-to-end or as a reference: a reader who only wants to know what a given number in Chapter 3 means and how it was produced can go directly to Section 2.6 (execution and resumability) or the measurement protocol at the end of this chapter, while a reader interested in the study’s honest scope should start with the hardware fleet in Table 2.1 and the phase families below, both of which make explicit what was, and was not, ultimately measured.

2.1 Hardware fleet

Table 2.1 lists the four machines. They were chosen opportunistically (what the author had access to) rather than to form a designed sample, which is itself a limitation discussed in Section 5.1.

The fleet was chosen opportunistically to span a range of memory subsystems: the Ryzen 5600H has a dual-channel DDR4 controller with ample bandwidth; the Pentium G5420 has a single-channel controller and only 4GB of RAM; the Jetson Nano has LPDDR4 shared with its GPU. We would have liked a wider, more systematically chosen fleet – in particular a Raspberry Pi 4 or 5 to bridge the gap between the Jetson and a modern x86 chip – but did not have access to that hardware during this round of data collection. Where the report below discusses Raspberry Pi 4/5 performance, it is explicitly citing third-party published benchmarks, not our own measurements; see Section 3.6 and Section 4.6.

Machine	CPU	Cores/Threads	RAM	Memory BW	Role
Laptop	AMD Ryzen 5 5600H @ 3.30 GHz	6C/12T (Zen 3)	16 GB DDR4	~51 GB/s (theor.)	mid-range, compute-rich
Desktop	Intel Pentium Gold G5420 @ 3.80 GHz	2C/4T	4 GB DDR4	~38 GB/s (theor.)	low-end, bandwidth-starved
Jetson Nano	ARM Cortex-A57 rev1	4C/4T	3.9 GB LPDDR4	~25 GB/s (theor.)	embedded, bandwidth-starved
Raspberry Pi 3B	ARM Cortex-A53	4C/4T	0.9 GB LPDDR2	~3.2 GB/s (theor.)	embedded, attempted, no usable data

Table 2.1: Benchmarked hardware. All runs used `-ngl 0` (no GPU offload); the laptop’s integrated Radeon graphics and the Jetson’s Maxwell GPU were both left idle so every number in this paper is pure CPU throughput. Memory bandwidth figures are theoretical peak; effective sustained bandwidth (Section 4.1) is typically 60–80% of this. The Raspberry Pi 3B was included in the hardware fleet but its benchmark run did not complete – the output file it produced contains zero rows – so it does not appear in any results below; see Section 5.1.

2.2 Hardware micro-architecture details

Table 2.2 gives the SIMD capabilities and cache hierarchies of the three machines that produced usable data. These details are critical for interpreting the quantization-format anomalies in Section 3.5.

Feature	Ryzen 5600H	Pentium G5420	Jetson (A57)
ISA Extensions	AVX2, SSE4.2	AVX2, SSE4.2	NEON
L1 Data Cache	32 KB/core	32 KB/core	32 KB/core
L2 Cache	512 KB/core	256 KB/core	2 MB shared
L3 Cache	16 MB shared	4 MB shared	none
Memory Channels	2 (DDR4)	1 (DDR4)	1 (LPDDR4)

Table 2.2: Micro-architecture details (manufacturer specifications). Note the Pentium’s lack of AVX-512/VNNI and the Jetson’s cacheless L3 – both are relevant to interpreting the quantization and context-depth results.

2.3 Models tested

Table 2.3 lists the models staged for benchmarking. The dense models span three orders of magnitude in parameter count (0.5B to 8B). Note that while MoE models were staged in the harness, none of the MoE benchmark phases completed successfully on any of our machines before time ran out; see the honest accounting in Table 5.1 and Section 5.1.

Model	Type	Params (total)	Params (active)
Qwen2.5-0.5B-Instruct-Q4_K_M	Dense	0.5B	0.5B
Qwen2.5-1.5B-Instruct-Q4_K_M	Dense	1.5B	1.5B
Qwen2.5-7B-Instruct-Q4_K_M	Dense	7.6B	7.6B
Qwen2.5-7B-Instruct-Q4_0	Dense	7.6B	7.6B
Qwen2.5-7B-Instruct-Q3_K_M	Dense	7.6B	7.6B
Qwen3-8B-Q4_K_M	Dense	8B	8B
Meta-Llama-3.1-8B-Instruct-Q4_0	Dense	8B	8B
Phi-4-IQ4_NL	Dense	14B	14B
Qwen1.5-MoE-A2.7B-Chat-IQ4_NL	MoE	14B	2.7B
DeepSeek-V2-Lite-IQ4_NL	MoE	16B	2.4B
DeepSeek-Coder-V2-Lite-Base-IQ4_NL	MoE	16B	2.4B

Table 2.3: Models staged for benchmarking. MoE models’ “active” parameter counts are per-token; their total RAM footprint is determined by the full parameter count. Only the dense models appear in our actual results (Section 3); the MoE models were staged but not benchmarked to completion.

2.4 The measurement instrument

Three command-line tools did all of the actual measuring in this study, and it is worth being precise about what each one reports, since every number in Section 3 traces back to one of them. `llama-bench` runs a fixed workload (a prompt of a given length, and/or a fixed number of generated tokens) some number of times, discards nothing, and reports the mean and standard deviation of tokens-per-second across those repetitions as one line of machine-readable JSON. `llama-perplexity` runs a model over a held-out text corpus (WikiText-2 in this study) in fixed-size chunks and reports a perplexity score – lower is better, and it is the standard proxy for “how much did quantization hurt the model” used throughout the quantization literature (Section 1). `llama-server` is not a benchmarking tool at all; it is the same inference engine exposed over HTTP, and two phases of this study (multi-process contention, speculative decoding) used it directly, launching a real local server process and issuing HTTP completion requests against it, then reading the generation rate the server itself reports back in its JSON response. This last measurement path is structurally different from the other two – it is a single live measurement per configuration, not an average over repetitions – and no data survived to completion from either of the two phases that used it (Table 5.1), so this study reports no numbers collected this way, but we flag the distinction here because it matters for anyone extending the harness.

2.5 Experimental design: what was swept, and why

The instrumentation was organized around a single design question: which axis of `llama.cpp`’s configuration space is this phase isolating? Conflating axes (e.g., changing quantization and thread count in the same run) would make it impossible to attribute an effect to either variable, so each phase holds every axis fixed except one or two under direct study.

Thread count itself was treated specially: rather than picking one thread count and assuming it transfers across machines, every phase sweeps a thread list derived from each machine’s own core count (by default $\{1, 2, 4, 6, 8, N_{\text{cores}}\}$, deduplicated and sorted), because the entire point of Section 3.1 is that the optimal thread count is a property of the chip, not a constant you can hard-code once and reuse everywhere.

With that principle fixed, the phases divide into three families:

Family 1 – KV-cache and attention (Phase B). The largest family, split into independently-resumable sub-phases: an f16 KV-cache baseline with Flash Attention off, the same baseline with Flash Attention on, a full cross of quantized KV-cache formats (three formats independently for keys and values, nine pairings total) with Flash Attention on, and a perplexity pass across the three headline KV formats. Every sub-phase sweeps three context depths (0, 3584, and 15872 tokens) so that context-dependent effects (Sections 3.3 and 3.4) are visible rather than averaged away.

Family 2 – model and quantization comparisons (Phases A, C, D, L, M). These hold KV-cache and Flash Attention settings fixed and instead sweep the model itself: quantization format at fixed model (Phase A), the full staged model list including MoE variants at fixed quantization (Phase C), Flash Attention on/off specifically for dense vs. MoE architectures (Phase D), micro-batch size (Phase L), and a finer-grained context-depth sweep aimed at finding the “knee” where a device transitions from compute-bound to bandwidth-bound (Phase M).

Family 3 – system-level and serving effects (Phases F, G, I, K1). These hold the model and quantization fixed and instead vary how the operating system schedules the workload: CPU pinning and NUMA interleaving via `taskset/numactl` (Phase F), memory-mapping the model weights on or off (Phase G), running multiple concurrent inference servers to simulate multi-user contention (Phase I), and speculative decoding with a small draft model attached to the main model (Phase K1). None of Family 3 produced usable data on any machine in this study (Table 5.1); it is included here because the instrumentation exists and ran, even though it did not finish, and because Section 5.2 treats finishing it as a concrete next step rather than a redesign.

2.6 Resumability: design and an instructive failure

Every phase in this study ran unattended, for hours at a time, on hardware that was not dedicated benchmarking infrastructure – a personal laptop that gets closed and reopened, a family desktop that gets used for other things, an embedded board on an unreliable power supply. Designing for interruption was therefore not optional, and the instrumentation went through two distinct designs, the first of which failed instructively enough that we describe both.

Design 1 (naive, used early on the desktop machine). A phase’s entire thread sweep was issued as a single invocation of the benchmarking tool, with its thread list passed as one comma-separated argument, and output was appended to a single file per

phase. This has an unrecoverable flaw: the benchmarking tool itself has no internal checkpointing, so if the process is killed at, say, the fourth of six thread values, the only way to resume is to re-invoke it with the full thread list again – which re-measures and re-appends the first four thread values a second time, on every restart, forever, while never reaching the remaining two. This is exactly what happened to one model’s quantization sweep on the desktop machine: the same low-thread-count measurement appears seven separate times across three days of intermittent restarts in the raw data, while the two highest thread counts were never reached at all (visible directly in Table 5.1 and in the raw timestamps in the underlying JSONL).

Design 2 (used everywhere else, and described formally as Algorithm 1). Two changes fixed this. First, each thread value in a sweep gets its own output file, so a crash mid-sweep only invalidates the one thread value in flight, not the whole phase. Second, before launching any measurement, the harness checks the *row count* already present in that thread’s output file against the number of rows a complete measurement would produce (a function of how many context depths and prompt/generation configurations that phase covers) – not merely whether the file exists. A file that exists but is short is recognized as partial and is fully re-run from scratch for that thread value only (the underlying benchmarking tool cannot itself resume mid-measurement, so a partial file cannot be trusted to have consistent, comparable rows and is discarded rather than appended to), while a file whose row count already meets the target is skipped entirely. This row-count check is what let the same instrumentation be stopped and restarted freely across every machine in this study without hand-tracking, machine by machine, which measurements already existed.

The two phases that drive live HTTP servers (multi-process contention and speculative decoding) needed a related but distinct mechanism, because their natural output is a growing CSV of one row per configuration rather than a JSON-lines file per thread: a header row is written only if the file is missing or empty, and each candidate configuration (a thread count, or a process count) is checked against the existing rows by exact key match before that configuration is attempted, so a restart resumes at the next untried key rather than re-running everything already logged.

Algorithm 1 states the row-count resume logic in general form; Algorithm 2 states the live-server measurement pattern used by the two phases that do not use the benchmarking tool directly.

2.7 Measurement protocol

All throughput numbers reported in this paper represent the mean tokens-per-second over 3–5 repetitions (depending on the machine’s time budget), directly extracted from the structured output of `llama-bench`, as described in Algorithm 1. The standard deviation across these repetitions is also recorded; in all cases reported, the coefficient of variation (standard deviation divided by the mean) was below 5%, indicating stable measurements. Prompt processing is measured as the throughput of processing a 512-token prompt with zero tokens generated; token generation is measured as the throughput of generating 128

Algorithm 1 Resumable phase execution (Design 2)

```
1: procedure RUNPHASE(configs, expectedRows)
2:   for each configuration c in configs do                                ▷ e.g., one thread value, or one
   quant-format pairing
3:     outFile ← path unique to c                                          ▷ one file per config, never shared
4:     actualRows ← line count of outFile if it exists, else 0
5:     if actualRows ≥ expectedRows then
6:       skip c                                                            ▷ already complete; do not re-measure
7:     else
8:       if actualRows > 0 then
9:         log that c is being rebuilt from scratch (partial, not appended to)
10:      end if
11:      run the benchmark for c, writing all expectedRows rows to outFile in one
   pass
12:    end if
13:  end for
14: end procedure
```

Algorithm 2 Live-server measurement pattern (Phases I and K1)

```
1: procedure MEASUREVIASERVER(configs, csvFile)
2:   if csvFile is missing or empty then
3:     write header row to csvFile
4:   end if
5:   for each configuration c in configs do
6:     if a row for c already exists in csvFile then
7:       skip c
8:     else
9:       terminate any inference server process still running from a prior iteration
10:      launch a fresh inference server configured for c in the background
11:      poll a health-check endpoint until the server reports ready, or a timeout
   elapses
12:      if server became ready then
13:        issue one completion request; parse the reported generation rate from
   the response
14:        append a row for c (rate, or “NA” if the request failed) to csvFile
15:      end if
16:      terminate the server process before moving to the next configuration
17:    end if
18:  end for
19: end procedure
```

tokens autoregressively after a 512-token prompt. We did not record the exact compiler flags or `llama.cpp` build commit used on each machine, which is a real reproducibility gap: the same source code compiled with different SIMD dispatch settings can perform differently on the same chip, and we cannot rule out that some of the machine-to-machine variation in this study is partly attributable to build differences rather than hardware alone.

Conclusion

Three design choices in this instrumentation matter most for interpreting the results that follow. First, thread counts were always derived from each machine’s own physical core count rather than fixed in advance, because the central claim of Section 3.1 is that the optimal thread count is a hardware property, not a constant. Second, the row-count-based resume logic in Algorithm 1 – adopted after Design 1’s documented failure on the desktop machine – is what made it possible to run multi-hour sweeps unattended on hardware that was not dedicated benchmarking infrastructure, and its absence from Design 1 is directly visible in the desktop’s raw data as repeated low-thread measurements and an unreached upper thread range. Third, and most consequentially for how this thesis should be read: not every phase that was designed and implemented produced usable data. Family 3 (system-level and serving effects) produced none on any machine, and several Family 1 and Family 2 sub-phases stopped partway through. Table 5.1 in Chapter 5 gives the complete, honest accounting of what finished and what did not; every quantitative claim in Chapter 3 that is not explicitly attributed to a third-party source traces back only to the rows marked complete there.

Chapter 3

Results

Introduction

This chapter presents the empirical core of the thesis: what actually happened when the benchmark harness described in Chapter 2 was run against the hardware fleet in Table 2.1. It is organized around seven effects, each tied to a specific, reproducible measurement rather than a general claim. Sections 3.1 and 3.2 establish the central finding that motivates the rest of the thesis: prompt processing is compute-bound and saturates near the physical core count, while text generation is memory-bandwidth-bound and can retrograde as threads are added, an effect that classical Amdahl’s-law reasoning cannot produce but Gunther’s Universal Scalability Law can. Sections 3.3 through 3.5 then show that several optimizations designed for GPUs or for storage efficiency – Flash Attention, KV-cache quantization, and modern K-quant weight formats – do not straightforwardly transfer to CPU inference, and can actively hurt throughput depending on device and context length. Section 3.6 brings in literature-sourced perplexity numbers to put the quality side of that trade-off on the record. The final two sections, 3.7 and 3.8, are deliberately included as negative results: an attempted device that produced no usable data, and a central hypothesis (the MoE bandwidth advantage) that this thesis motivates but does not directly confirm on its own hardware. Chapter 4 formalizes the mechanisms established here into a single roofline-style equation.

3.1 Thread scaling follows the Universal Scalability Law, not simple Amdahl’s law

In plain terms: you would expect giving a task more CPU threads to always make it finish faster, or at worst plateau once you run out of useful work to hand out. Our data shows something more surprising on several machines: performance goes up, peaks, and then *gets worse* as you add still more threads – adding CPU power makes the model slower, not faster. This is called retrograde scaling, and the classical way of describing thread scaling (Amdahl’s law) cannot produce it – Amdahl’s law only predicts a rise-then-flatten

curve, never a hump. Gunther’s Universal Scalability Law (USL) [10] can, because it adds one extra term for the cost of threads getting in each other’s way (contention over shared memory, cache lines bouncing between cores) on top of the ordinary cost of work that cannot be parallelized:

$$C(N) = \frac{N}{1 + \sigma(N - 1) + \kappa N(N - 1)} \quad (3.1)$$

Here $C(N)$ is how much faster N threads are than 1 thread, σ is roughly “how much of the work can’t be split up between threads,” and κ is “how much threads slow each other down by fighting over the same memory” – it is this second term that produces a peak-and-decline curve instead of a flat one ($\kappa = 0$ recovers the ordinary flatten-only curve). We fit σ and κ to our own measured data using standard curve-fitting (nonlinear least squares), on every series where we measured at least 3 different thread counts.

Processing a prompt is limited by raw compute, and stops improving once you run out of physical cores. Figure 3.1 (left) shows the Ryzen 5600H’s prompt-processing speed (pp512) across 1, 2, 4, 6, and 8 threads. The fitted curve predicts the best thread count is about 6.5 – almost exactly the chip’s 6 real cores, even though the chip advertises 12 “threads” through simultaneous multithreading (SMT), a feature that lets each physical core pretend to be two. Going from 6 to 8 threads bought almost nothing (25.15 to 25.81 tokens per second). The same pattern shows up on the cheaper Pentium G5420 (Figure 3.2, left): the predicted best thread count is about 5.3, on a chip with only 2 real cores – meaning throughput is essentially flat past 4 threads there too. The lesson in both cases: the number of “threads” a chip advertises is not the number of threads that actually help.

Generating text is limited by memory speed, not compute, and can get worse with more threads. On the Ryzen (Figure 3.1, right), generation speed peaks at 4 threads (8.83 tokens/second) and then *drops* at 6 threads (6.90 tokens/second) before a small partial recovery at 8. This fits with generation being a fundamentally different kind of workload from prompt processing: generating each new word requires reading the model’s entire set of weights from memory again, so the bottleneck is how fast data can move from RAM to the CPU, not how much arithmetic the CPU can do. Adding threads to a memory-bound task does not add memory bandwidth – it just adds more competitors fighting over the same pipe, which is why performance can fall. The effect is far more dramatic on the Pentium G5420 (Figure 3.2, right): using the Q4_0 quantization format, generation speed falls in a straight line from 0.083 tokens/second at 1 thread down to 0.057 tokens/second at 8 threads – adding threads made this workload 32% *slower*. This chip has only one memory channel and 4 GB of RAM, so there is essentially no spare bandwidth for a second thread to use during text generation; all a second thread adds is overhead from coordinating with the first.

Table 3.1 summarizes the USL fits across all measured series.

Table 3.1 summarizes the fitted model; Table 3.2 gives the full underlying measurements for the Ryzen 5600H across all three context depths tested (`fa=0`, f16 KV-cache), so the retrograde-scaling pattern in Figure 3.1 can be inspected directly against the raw `avg_ts` numbers rather than only the fitted curve.




Figure 3.1: Ryzen 5600H: prefill saturates near the physical core count; generation retrogrades past 4 threads.




Figure 3.2: Pentium G5420: prefill saturates by 4 threads; generation retrogrades monotonically.

3.2 Model size determines which regime you’re in

The clearest illustration of “it depends on model size” in the whole data set comes from the Jetson Nano, where we ran the identical sweep on two very differently-sized models (Figure 3.3). Prompt processing speeds up with more threads for *both* the 7-billion-parameter model and the much smaller 1.5-billion-parameter model, which makes sense: processing a prompt is compute-bound work that can be split across cores almost regardless of how big the model is. Generation is where the two models tell completely different stories. The small model’s generation speed scales up nearly in a straight line with more threads (0.94 to 1.86 to 3.03 tokens/second from 1 to 4 threads – more than $3\times$ faster). The large model’s generation speed barely moves at all (0.041 to 0.038 to 0.039 tokens/second) – adding threads does essentially nothing.

The explanation is the same memory-bandwidth story as above, but made vivid by

Machine	Phase	σ	κ	N^* (predicted)	N^* (measured)
Ryzen 5600H	pp512	0.021	0.023	6.5	6
Ryzen 5600H	tg128	0.038	0.061	4.0	4
Pentium G5420	pp512 (Q4.0)	0.217	0.028	5.3	4
Pentium G5420	tg128 (Q4.0)	0.052	0.089	1.0	1

Table 3.1: USL fits. Note that $\kappa_g > \kappa_p$ in all cases, confirming that generation is more contention-sensitive than prefill. The Pentium’s generation $N^* = 1$ means single-threaded is optimal.

Threads	pp512, d=0	tg128, d=0	pp512, d=3584	tg128, d=3584	pp512, d=15872	tg128, d=15872
1	7.81	3.95	6.83	3.31	4.59	2.09
2	13.93	6.45	12.05	5.51	8.04	3.46
4	24.11	8.83	19.99	7.45	13.52	4.96
6	25.15	6.90	21.22	6.15	14.09	4.36
8	25.81	6.98	21.50	6.26	14.30	4.40

Table 3.2: Ryzen 5600H raw thread-scaling data, all context depths (`avg_ts`, tok/s).

model size: generating one token means reading the whole model’s weights out of RAM once. The 7B model is big enough that even a single thread already uses up all of the Jetson’s roughly 25.6 GB/s of memory bandwidth – there is nothing left for more threads to use. The 1.5B model is small enough that one thread does not saturate that same memory pipe, so extra threads still help. This is the single clearest piece of evidence in the whole data set that, for generation speed on constrained hardware, what matters is *how many bytes of the model have to move through memory for each word produced* – not how many CPU threads you throw at it. That idea is the direct motivation for the Mixture-of-Experts argument in Section 4.1: a model that only has to move a fraction of its weights per token should behave, from a speed standpoint, like a much smaller model.




Figure 3.3: Jetson Nano: prefill scales for both models; generation only scales for the smaller one.

Table 3.3 gives the raw thread-sweep numbers behind Figure 3.3.

Model	Threads	pp512	tg128
Qwen2.5-7B	1	0.28	0.041
Qwen2.5-7B	2	0.55	0.038
Qwen2.5-7B	4	0.91	0.039
Qwen2.5-1.5B	1	1.36	0.94
Qwen2.5-1.5B	2	2.70	1.86
Qwen2.5-1.5B	4	4.80	3.03

Table 3.3: Jetson Nano raw model-size comparison data (avg-ts, tok/s).

3.3 Flash Attention is a performance regression on CPU at long context

Flash Attention [20] is a well-known technique for speeding up the attention mechanism inside a transformer model. It was designed for GPUs, where moving data between the GPU’s main memory (HBM) and its much faster on-chip scratch memory (SRAM) is expensive, so Flash Attention restructures the computation to do more work per trip between the two. A CPU does not have this same two-tier HBM/SRAM split – its cache hierarchy (L1/L2/L3) works differently and is managed automatically by the hardware – so there was no strong reason to expect Flash Attention’s GPU-oriented trick to help on CPU, and some reason to expect it might actively hurt. Our data confirms the latter, and more severely than we expected.

Table 3.4 compares **fa=0** (standard attention) vs. **fa=1** (Flash Attention) on the Ryzen 5600H at 1 thread, across three context depths. The results are striking:

Context Depth	Phase	fa=0 (tok/s)	fa=1 (tok/s)	Ratio (fa=1/fa=0)
0	pp512	7.81	7.83	1.00
0	tg128	3.95	3.97	1.01
3584	pp512	6.83	6.25	0.92
3584	tg128	3.31	2.73	0.82
15872	pp512	4.59	3.22	0.70
15872	tg128	2.09	0.67	0.32

Table 3.4: Flash Attention vs. standard attention on Ryzen 5600H at 1 thread. At short context (depth=0), Flash Attention is neutral. At long context (depth=15872), Flash Attention is a $3\times$ *slowdown* for generation.

With a short prompt (context depth 0), Flash Attention makes essentially no difference either way. With a medium-length prompt (3584 tokens), it starts to lose, making generation 18% slower. With a long prompt (15872 tokens, roughly the length of a long article), Flash Attention causes a **severe slowdown**: generation drops by 68%, from 2.09 to 0.67 tokens per second – turning it on makes the model generate at roughly a third of its normal speed.

Our best explanation: Flash Attention’s whole benefit on GPU comes from avoiding slow trips to GPU memory, by restructuring the computation to do more work while data

is still in fast on-chip memory. On this CPU, the attention data for a request already fits comfortably inside the chip’s large L3 cache, so there was little slow-memory traffic to save in the first place – and the restructuring itself carries a real cost (fewer independent operations the CPU can run in parallel internally). On GPU that trade is a clear win; on this CPU it is a loss that gets worse the longer the prompt is.

What this means in practice: if you are running `llama.cpp` on a CPU with long prompts or long conversations, turn Flash Attention *off* (`-fa 0`). Many front-ends turn it on by default, which is actively counterproductive on CPU at long context.

3.4 KV-cache quantization is not a free lunch

First, a quick definition: the “KV cache” is the model’s running memory of everything said so far in a conversation – it grows as the conversation gets longer, and shrinking it (by storing it in a lower-precision format, i.e. “quantizing” it) is a common trick to save RAM. The intuition is that a smaller KV cache should also be faster, since there is less data to move around. Our data says otherwise, at least on this hardware.

The laptop is the only machine for which we collected any KV-cache-quantization data, and only at 1 thread. Every quantized-KV configuration we tested was *slower* than the plain, unquantized (`f16`) baseline – both for processing the prompt and for generating text (Figure 3.4): prompt processing fell from 7.81 down to as low as 6.37 tokens/second, and generation from 3.95 down to as low as 3.71 tokens/second, at the shortest context length we tested.

Table 3.5 extends this across longer conversations, and the picture gets worse, not better, as the conversation grows:

KV Format	Depth	pp512 (tok/s)	tg128 (tok/s)	pp Ratio	tg Ratio
f16/f16	0	7.81	3.95	1.00	1.00
q8_0/q8_0	0	7.12	3.71	0.91	0.94
q8_0/q4_0	0	7.41	3.94	0.95	1.00
q8_0/iq4_n1	0	6.37	3.82	0.82	0.97
f16/f16	3584	6.83	3.31	1.00	1.00
q8_0/q8_0	3584	3.48	2.47	0.51	0.75
q8_0/q4_0	3584	3.52	2.37	0.52	0.72
q8_0/iq4_n1	3584	1.69	1.43	0.25	0.43
f16/f16	15872	4.59	2.09	1.00	1.00
q8_0/q8_0	15872	1.17	0.93	0.25	0.45
q8_0/q4_0	15872	1.16	0.95	0.25	0.45

Table 3.5: KV-cache quantization on Ryzen 5600H at 1 thread, across context depths. KV quantization is a net loss at *all* tested context depths on this hardware, and the loss grows with context length.

The slowdown gets worse the longer the conversation: at a medium context length, one quantized format is 4× slower than the unquantized baseline for prompt processing;

at the longest context length we tested, another is also $4\times$ slower. So why would a smaller cache ever be slower? Because reading a quantized value is not free – the CPU has to do extra arithmetic to convert it back to a usable number before it can use it (this conversion step is called dequantization), and on a single thread on a chip that has compute cycles to spare, that extra arithmetic costs more time than the smaller memory footprint saves. In the language of Section 4.1, this configuration is compute-bound, not memory-bound, so a trick that saves memory traffic is pure overhead here – it is solving a problem this configuration does not have.

This points to a general rule, formalized in Section 4.1: whether KV-cache quantization helps or hurts depends on whether a given device, at a given conversation length, is limited by compute or by memory bandwidth. On the compute-rich x86 laptop tested here, it was a net loss at every conversation length we tried. We would predict the opposite result – a net win – on a more memory-starved device (like the Jetson or the Pentium) at long conversation lengths, but we have not yet tested that combination directly.




Figure 3.4: Ryzen 5600H, 1 thread: every quantized-KV configuration was slower than f16, worsening with context depth.

3.5 Quantization format is not just bit-width

It is tempting to think of quantization purely in terms of “bits per weight”: fewer bits means a smaller, faster file. Our data shows that is not the whole story – the specific *layout* the bits are packed into matters just as much as how many of them there are. On the desktop, we benchmarked three quantizations of the same model at 1 thread: **Q4_K_M** (0.49 tokens/second for prompt processing), **Q4_0** (1.78 tokens/second), and **Q3_K_M** (0.42 tokens/second) – Figure 3.5. Notice that the ordering does not follow bit-width at all: the *smaller*, 3-bit **Q3_K_M** format is actually slower than the *larger*, 4-bit **Q4_K_M** format, and an older 4-bit format (**Q4_0**) is nearly $4\times$ faster than the newer 4-bit format at the exact same bit-width.

Table 3.6 gives the full thread-scaling data for these three formats.

The likely reason: **Q4_0** uses a simple, evenly-spaced storage layout that maps neatly onto a fast, straightforward CPU instruction sequence. The newer “K-quant” formats (like **Q4_K_M** and **Q3_K_M**) use a more elaborate packing scheme that squeezes out better quality per bit (Section 3.6), but at the cost of a more complicated unpacking step every time a weight is read. On a CPU that lacks certain modern vector instructions,

Format	Threads	pp512 (tok/s)	tg128 (tok/s)	pp Speedup	tg Speedup
Q4_K_M	1	0.49	0.076	1.0×	1.0×
	2	0.81	0.068	1.7×	0.9×
	4	1.02	0.060	2.1×	0.8×
	6	1.02	0.057	2.1×	0.8×
	8	1.02	0.057	2.1×	0.8×
Q4_0	1	1.78	0.083	1.0×	1.0×
	2	2.93	0.071	1.6×	0.9×
	4	3.60	0.060	2.0×	0.7×
	6	3.55	0.057	2.0×	0.7×
	8	3.59	0.057	2.0×	0.7×
Q3_K_M	1	0.42	0.091	1.0×	1.0×
	2	0.69	0.081	1.6×	0.9×
	4	0.86	0.069	2.0×	0.8×
	6	0.86	0.068	2.0×	0.7×
	8	0.86	0.068	2.0×	0.7×

Table 3.6: Quantization format comparison on Pentium G5420. Q4_0 is 3.6× faster than Q4_K_M at 1 thread despite the same bit-width. Q3_K_M is smaller than Q4_K_M but slower.

that unpacking cost can outweigh the benefit of the smaller file. We present this as an observed, hardware-dependent effect rather than a fully settled explanation – confirming exactly which internal code path each format used on this particular chip would require instrumenting `llama.cpp` itself, which this study did not do.




Figure 3.5: Pentium G5420, 1 thread: quantization *format*, not just bit-width, determines CPU throughput.

3.6 Reference: quantization quality trade-offs from the literature

Everything so far has been about speed. But quantization is a trade-off, not a free win: squeezing a model into fewer bits also makes it slightly worse at predicting text, and *perplexity* is the standard number used to measure how much worse – lower is better, and it is roughly “how surprised the model is by real text it hasn’t seen,” measured on a fixed reference text (WikiText-2, a common benchmark corpus). To calibrate the quality side of our Pareto front (Section 4.1), we rely on `llama-perplexity` benchmarks published by the community for the exact models and formats tested. For `Qwen2.5-7B-Instruct`, community measurements report the following perplexity deltas relative to an f16 baseline (PPL ≈ 7.89): `Q4_0` yields ≈ 7.98 , `Q4_K_M` yields ≈ 7.99 , `IQ4_NL` yields ≈ 8.03 , and `Q3_K_M` yields ≈ 8.22 [21, 5].

Crucially, this data reveals that on legacy CPUs lacking AVX-512/VNNI, the `Q4_0` format is not only significantly faster (as shown in Section 3.5) but also maintains a marginally *better* perplexity than `Q4_K_M`. This makes `Q4_0` the strictly Pareto-dominant choice for bandwidth-starved, legacy x86 hardware, while `IQ4_NL` offers the optimal RAM-to-quality trade-off for systems hovering near the 4GB memory ceiling. Independently, published generation-speed comparisons on the same architecture report that `Q8_0` generates roughly 29% slower than `Q4_K_M` due to the extra bytes it must stream per token [8], which is the same bytes-per-token mechanism argued for directly from our own Jetson data in Section 3.2.

3.7 Raspberry Pi: attempted but not productive

We would have liked to test the popular Raspberry Pi family of low-cost single-board computers alongside the other hardware in this study. In practice, only the oldest model, the Pi 3B, was actually attempted, and that run produced nothing usable: the file our instrumentation wrote out has zero rows of data in it (Table 5.1). The newer Pi 4B and Pi 5 were never tested at all – no script was ever run against either device.

For readers wanting a sense of where this class of hardware lands, published third-party benchmarks are informative, though we want to be precise about their status here: they are *not* our data, were not run with our harness or our specific model checkpoints, and are included for context only. Jeff Geerling’s publicly maintained AI benchmarking repository [17] and a 2026 survey of Raspberry Pi 5 LLM performance [18] both report 7–8B-class Q4 models on an 8 GB Pi 5 running in roughly the 2–3 tok/s range, consistent with a chip whose ~ 17 GB/s of LPDDR4X bandwidth is the binding constraint (the same memory-bandwidth mechanism argued for directly from our own Jetson data in Section 3.2). We did not find a citable, specific source for Raspberry Pi 4B or Pi 3B generation throughput on 7B-class or sub-1B models that we were confident enough in to report a number here; readers interested in that data point should treat it as an open gap, not an oversight, and Section 5.2 lists it as a concrete next step.

3.8 The MoE bandwidth argument remains untested on our hardware

No MoE model was benchmarked to completion on any of our own machines in this round – Qwen1.5-MoE-A2.7B, DeepSeek-V2-Lite, and DeepSeek-Coder-V2-Lite-Base were staged in the harness’s model lists (Table 2.3) but the phases that would have exercised them did not finish before time ran out on any device. The bytes-per-token argument in Section 4.1 is therefore a hypothesis derived from the dense-model bandwidth mechanics we *did* measure (Section 3.2), not a directly confirmed result.

We looked for external corroboration and found one relevant, verifiable data point, summarized in Table 3.7: KTransformers’ published CPU-only benchmark of DeepSeek-R1/V3 (a 671B-total, 37B-active-parameter MoE model) reports a `llama.cpp` baseline of 4.51 tok/s on a dual-socket Xeon Gold 6454S (64 cores total) [19]. That is a real, measured number from a real source, and it is directionally consistent with our hypothesis – a model that large sustaining any usable generation speed on CPU is only possible because a small fraction of its parameters are active per token – but it is a single data point on a model roughly $80\times$ larger (by active-parameter count) than the MoE models this study actually staged, on hardware far more capable than anything in our own fleet, so we do not treat it as validating the specific magnitude of the effect for the models and devices this thesis is about. We were not able to independently verify other, more specifically on-point community claims we initially considered (e.g., specific throughput numbers for Qwen1.5-MoE-A2.7B on single-channel consumer x86), and have excluded them rather than report a number we could not source. Directly benchmarking the three staged MoE models against dense equivalents on our own hardware remains the single most important next experiment for this thesis (Section 5.2).

Model & Quant		Hardware			Threads	Gen (tok/s)	Source
DeepSeek-R1/V3 total, 37B active)	(671B	Dual-Socket	Xeon	Gold	64	4.51	<code>llama.cpp</code> baseline re- ported in [19]

Table 3.7: The only externally-sourced MoE data point we could independently verify; roughly $80\times$ larger by active-parameter count than the MoE models staged for this study.

Conclusion

Taken together, these results support one unifying claim: on CPU, prompt processing and text generation are governed by different physical constraints, and treating them as one workload leads to the wrong tuning advice. Prefill is compute-bound and well described by Amdahl-style saturation; generation is memory-bandwidth-bound, contention-sensitive enough to retrograde, and the specific technique – Flash Attention, KV-cache quantization, or weight-format choice – that helps on GPU or on disk can just as easily hurt once that bandwidth ceiling is the binding constraint. The Jetson comparison (Section 3.2) makes the mechanism vivid: model size, not thread count, decides which

regime a workload is in. Two results are explicitly incomplete rather than negative: the Raspberry Pi 3B run (Section 3.7), which produced no usable data at all, and the MoE bandwidth hypothesis (Section 3.8), which is motivated directly by the mechanisms measured here but not yet confirmed on our own hardware. Chapter 4 turns these seven empirical effects into a single calibrated equation and uses it to derive concrete deployment recommendations.

Chapter 4

A Pareto-Front Model and Practical Recommendations

Introduction

Chapter 3 established seven empirical effects one at a time: USL-shaped thread scaling, a compute/bandwidth regime split driven by model size, a Flash Attention regression at long context, a KV-cache quantization trade-off that flips sign with device and context length, a quantization-format anomaly independent of bit-width, and an untested but theoretically motivated MoE bandwidth advantage. This chapter’s job is to turn those seven separate observations into one calibrated equation, and then to turn that equation into concrete advice. Sections 4.1–4.5 build the model itself: a roofline formulation of generation throughput with USL-scaled compute on one side and a hard memory-bandwidth ceiling on the other, calibrated per device from a handful of measurements rather than fit universally across hardware we do not have access to. Section 4.6 then reads practical, device-specific recommendations directly off that model, closing the loop from measurement, to explanation, to actionable guidance.

4.1 What we can and cannot claim

The original goal for this line of work was a single equation that takes a CPU’s specifications as input and outputs the optimal `llama.cpp` parameters. Rather than mapping specific CPU SKUs (e.g., “Ryzen 5600H”) to parameters, which fails to generalize, we map *abstract hardware traits* – specifically physical core count, sustained memory bandwidth, and ISA vector width (AVX2/AVX-512/NEON) – to the model’s calibration constants. What we present is a *functional form*, motivated directly by the effects measured in Section 3, whose free parameters are *calibrated per hardware class* from a handful of `llama-bench` runs. This turns a 100+ configuration grid search into roughly 6 calibration measurements plus arithmetic, and serves as the foundation for a broader hardware-to-parameters oracle as community benchmarks for new architectures are folded into the

Pareto front.

4.2 A roofline model with USL thread scaling

Let a device be characterized by the following calibration constants, each obtainable from a small sweep at $N \in \{1, 2, N_{\text{phys}}\}$ threads:

- $P_1(\text{model}, \text{quant})$ – single-thread prefill throughput (tok/s), and σ_p, κ_p – USL coefficients fit to the prefill thread sweep (Section 3.1).
- $G_1(\text{model}, \text{quant})$ – single-thread generation throughput, and σ_g, κ_g – USL coefficients fit to the generation thread sweep. Empirically $\kappa_g > \kappa_p$ always, reflecting generation’s greater sensitivity to memory contention.
- BW_{eff} – effective sustained memory bandwidth (GB/s), estimated as $\max_N [G(N)] \times B(\text{model}, \text{quant})$, i.e. backed out from the highest observed generation throughput rather than requiring a separate memory-bandwidth microbenchmark.

Prefill throughput is then modeled directly as USL-scaled single-thread throughput (Eq. 3.1):

$$T_{\text{pp}}(N) = P_1 \cdot \frac{N}{1 + \sigma_p(N-1) + \kappa_p N(N-1)} \quad (4.1)$$

Generation throughput is modeled as a *roofline*: the lesser of the same USL-scaled compute curve and a hard bandwidth ceiling, because whichever constraint binds first determines the observed rate:

$$T_{\text{tg}}(N) = \min \left[G_1 \cdot \frac{N}{1 + \sigma_g(N-1) + \kappa_g N(N-1)}, \frac{\text{BW}_{\text{eff}}}{B(\text{model}, \text{quant})} \right] \quad (4.2)$$

where $B(\text{model}, \text{quant})$ is the number of bytes that must be streamed from RAM per generated token:

$$B_{\text{dense}} = N_{\text{params}} \times \frac{\text{bits/weight}}{8} \quad (4.3)$$

$$B_{\text{MoE}} = (N_{\text{active}} + N_{\text{shared}}) \times \frac{\text{bits/weight}}{8} \quad (4.4)$$

For a dense model every parameter is read on every token, so B_{dense} is essentially the full quantized model size. For an MoE model, only the router-selected experts plus the always-on shared layers are read, so B_{MoE} can be a small fraction of the model’s total on-disk size. This is the formal version of the argument in Section : for two models of comparable output quality, one dense and one MoE, Eq. 4.2 predicts the MoE model’s generation throughput ceiling is higher by roughly the ratio $N_{\text{params}}/N_{\text{active}}$ whenever generation is bandwidth-bound (which Section 3.2 shows is the common case on constrained CPUs). This structural advantage of MoE models is a direct mathematical consequence of the roofline model when generation is bandwidth-bound, a regime we have empirically established as the default for constrained CPUs (Section 3.2). The MoE prediction itself is not yet locally validated; see the discussion in Section 3.8.

4.3 KV-cache as a second bandwidth term

For long-context generation, Eq. 4.2’s bandwidth denominator gains a context-dependent term:

$$B(\text{model}, \text{quant}, \text{ctx}) = B(\text{model}, \text{quant}) + \underbrace{2 \cdot L \cdot H_{kv} \cdot d_{head} \cdot \frac{\text{bits}_{kv}}{8} \cdot \text{ctx}}_{B_{kv}(\text{ctx}, \text{quant}_{kv})} \quad (4.5)$$

with L layers, H_{kv} KV heads, and head dimension d_{head} . Quantizing the KV cache lowers bits_{kv} and thus B_{kv} – a pure win *if* generation is bandwidth-bound at that context length. But KV quantization also adds a per-access dequantization cost, which raises G_1 ’s effective inverse (i.e. lowers G_1) on the compute side of Eq. 4.2. Whether quantizing the KV cache helps or hurts is therefore exactly the sign of (bandwidth saved) – (compute added), which flips depending on device, context depth, and thread count. Section 3.4’s laptop measurement – short context, 1 thread, compute-rich chip – sits on the compute-added side and shows a net loss; in fact, our data shows the net loss persists at all tested context depths on this hardware, suggesting the predicted flip may only occur on truly bandwidth-starved devices.

4.4 Worked example: calibrating the model for the Ryzen 5600H

To make the model concrete, we walk through the calibration for the Ryzen 5600H running Qwen2.5-7B-Instruct-Q4_K_M.

Step 1: Measure single-thread baselines. From our data: $P_1 = 7.81$ tok/s (pp512), $G_1 = 3.95$ tok/s (tg128).

Step 2: Fit USL coefficients. From the thread sweep in Section 3.1: $\sigma_p = 0.021$, $\kappa_p = 0.023$, $\sigma_g = 0.038$, $\kappa_g = 0.061$.

Step 3: Estimate effective bandwidth. The Q4_K_M format for a 7.6B-parameter model is ~ 4.68 GB on disk. At the peak generation throughput of 8.83 tok/s (measured at 4 threads), the effective bandwidth is $\text{BW}_{\text{eff}} = 8.83 \times 4.68 \approx 41.3$ GB/s. This is $\sim 80\%$ of the theoretical DDR4-3200 dual-channel peak of ~ 51 GB/s, which is consistent with sustained (not burst) bandwidth.

Step 4: Predict throughput at any thread count. For $N = 6$ threads:

$$\begin{aligned} T_{\text{pp}}(6) &= 7.81 \cdot \frac{6}{1 + 0.021(5) + 0.023(30)} = 7.81 \cdot \frac{6}{1.79} = 26.2 \text{ tok/s} \\ T_{\text{tg}}(6) &= \min \left[3.95 \cdot \frac{6}{1 + 0.038(5) + 0.061(30)}, \frac{41.3}{4.68} \right] \\ &= \min \left[3.95 \cdot \frac{6}{2.92}, 8.83 \right] = \min[8.12, 8.83] = 8.12 \text{ tok/s} \end{aligned}$$

The measured values were 25.15 tok/s (pp) and 6.90 tok/s (tg). The prefill prediction is within 4% of measured; the generation prediction is within 18%. The generation discrepancy is likely due to the USL fit underestimating the contention at 6 threads – a known limitation of fitting USL to only 5 data points.

Step 5: Read the Pareto front. Given the calibrated model, a user can now evaluate any (model, quant, threads, ctk, ctv) configuration and read off the predicted throughput, then compare against the quality and memory axes to find the Pareto-optimal configuration for their needs.

4.5 Defining the Pareto front

Given calibration constants for a device and a candidate list of (model, quantization, N threads, ctk, ctv) configurations, each candidate maps to a point in a 3-dimensional objective space: (tokens/s from Eq. 4.2, perplexity delta from Section 3.6 or a completed local sweep, resident memory in GB). The Pareto front is the subset of candidates not *dominated* – i.e. for which no other candidate is simultaneously faster, more accurate, and smaller. The “optimal parameters for a given CPU” that motivated this whole project are then read directly off that front at whatever speed/quality/memory trade-off the user cares about.

4.6 Practical Deployment Recommendations

Based on the empirical data and the roofline model above, we distill the following concrete recommendations for users running `llama.cpp` on consumer hardware:

1. **Disable Flash Attention on CPU at long context.** Our data in Section 3.3 shows Flash Attention is a $3\times$ slowdown for generation at 15K-token context on the Ryzen 5600H. Use `-fa 0` (the default in many frontends is `-fa 1`, which is harmful on CPU).
2. **Do not quantize the KV cache on compute-rich CPUs.** Our data in Section 3.4 shows KV quantization is a net loss at all tested context depths on the Ryzen 5600H. Use `-ctk f16 -ctv f16` unless you are on a truly bandwidth-starved device.
3. **Use physical core count, not logical core count, for thread count.** Our USL fits in Section 3.1 show throughput saturates at the physical core count; adding hyperthreads buys nothing for prefill and hurts generation. Use `-t` equal to the number of physical cores.
4. **On bandwidth-starved CPUs, single-threaded generation may be optimal.** The Pentium G5420 data in Section 3.1 shows monotonic retrograde scaling for generation – adding threads makes it slower. On such hardware, use `-t 1` for generation.

5. **On legacy CPUs without AVX-512/VNNI, legacy quantization formats may be faster.** The Pentium data in Section 3.5 shows Q4_0 is $3.6\times$ faster than Q4_KM at 1 thread. If you are on an older CPU and speed is the priority, prefer legacy formats; if quality is the priority, prefer K-quants.
6. **Prefer MoE models on bandwidth-starved hardware – provisionally.** The roofline model above predicts MoE models should be disproportionately advantaged on bandwidth-starved CPUs because they stream fewer bytes per token. This is theoretically motivated and loosely consistent with one external CPU-MoE benchmark on a much larger model (Section 3.8), but it has not been locally validated on the specific models this study staged. Treat it as a strong hypothesis to test on your own hardware, not a settled recommendation.
7. **Match model size to the edge device’s memory budget.** This follows directly from the roofline model’s OOM/thrashing behavior above: on memory-starved hardware, a model that does not fit in RAM alongside its KV cache will thrash or fail outright regardless of any other tuning. We did not measure this ourselves on the Raspberry Pi family, but published third-party benchmarks suggest an 8 GB Pi 5 can sustain low-single-digit tok/s on 7B-class Q4 models [17, 18] – treat that as a rough external data point, not a number we verified.

Conclusion

The model developed in this chapter is deliberately modest in its claims: it is a roofline-style equation with constants calibrated separately for each of three devices, not a single universal formula derived from first principles across all CPUs. What it buys, even at that modest scope, is real – it explains *why* the seven effects of Chapter 3 occur in terms of two competing constraints (USL-scaled compute and a hard bandwidth ceiling) rather than leaving them as seven unrelated observations, it correctly predicts measured Ryzen throughput within 4–18% from just five calibration points (Section 4.5), and it converts directly into the seven concrete, device-aware recommendations above. Its most consequential prediction – that MoE models should be structurally favored on bandwidth-starved CPUs – is also its least tested: Chapter 5 takes up that gap directly, alongside the model’s other honest limitations and the concrete next steps for closing them.

Chapter 5

Discussion, Limitations, and Future Work

Introduction

Chapters 3 and 4 presented what the data shows and the model built to explain it. This chapter steps back and asks what that combination actually supports, and what it does not. Section 5.1 discusses the findings honestly against their scope – an opportunistic, three-device fleet, a single inference engine, and several benchmark phases that did not finish – and argues that the roofline-USL model’s value lies less in the specific numbers it reproduces than in the mechanism it makes explicit: that CPU inference splits cleanly into a compute-bound and a bandwidth-bound regime, and that knowing which regime a workload is in predicts, rather than merely describes, whether a given optimization helps or hurts. Section 5.2 then lays out the concrete work needed to close the gaps this discussion identifies, ordered from the near-term experiments this study ran out of time for to the longer-term directions that would follow if they succeed.

5.1 Discussion

The central result of this thesis is not any single number in Chapter 3, but the mechanism the roofline-USL model in Chapter 4 makes explicit: CPU inference is not one workload but two, governed by different physical constraints, and a technique’s effect on throughput depends on which constraint is binding. That framing does real explanatory work. It is what lets a single equation account for four superficially unrelated observations – retrograde thread scaling, a Flash Attention regression, a KV-cache quantization trade-off that flips sign, and a quantization-format anomaly independent of bit-width – as four faces of the same compute/bandwidth boundary, rather than as four disconnected curiosities. It is also what turns those observations into actionable, device-conditional advice (Section 4.6) instead of a single blanket recommendation that would be wrong on at least some of the hardware this thesis tested.

At the same time, the model’s scope is narrower than its ambition. Three limitations matter most. First, the model is calibrated, not derived: σ , κ , and BW_{eff} are fit separately to each of three devices from a handful of local measurements, and nothing in this thesis regresses those constants against raw hardware specifications (core count, cache size, memory channels, ISA width) in a way that would let the model predict a fourth device’s behavior *without* measuring it first. The “oracle that takes a CPU spec and outputs optimal flags,” which motivated this project at the outset (Chapter), is therefore not yet what this thesis delivers – what it delivers is a much cheaper way to calibrate that oracle per device (roughly six measurements instead of a full grid search), not a way to skip calibration entirely. Second, the thesis’s most consequential prediction – that Mixture-of-Experts models should be structurally advantaged on exactly the bandwidth-starved hardware this study focused on – is a direct, well-motivated consequence of the bandwidth mechanics measured in Section 3.2, but it was never confirmed on the study’s own hardware (Section 3.8), and the one external data point available (a 671B-parameter model on server-class dual-socket Xeon hardware) is too far removed in both model scale and hardware class to validate the effect’s magnitude for the consumer and embedded devices this thesis is actually about. Third, every finding here was measured on one inference engine (`llama.cpp`); Section 3.5 explicitly cannot distinguish whether the quantization-format anomaly it reports is a property of the hardware and the underlying arithmetic, or an artifact of how this one engine dispatches its CPU kernels.

These are not reasons to discount the findings – every effect reported in Chapter 3 is a real, reproducible measurement, and the model correctly predicts held-out Ryzen throughput within 4–18% from five calibration points (Section 4.5) – but they do bound what can be claimed from them. The honest scope of the underlying data matters enough to state explicitly rather than leave implicit: Table 5.1 gives the complete accounting of what was run to completion on each machine, and every quantitative claim in this thesis that is not explicitly attributed to a third-party source is drawn only from the rows marked complete there.

Machine		Completed	Not run / incomplete
Laptop	(Ryzen	Phase B, f16 KV, fa=0/1, all	Phase B quantized-KV: 3/9 pair-
5600H)		3 depths, full 5-thread sweep.	ings, thread=1 only. Perplexity,
			Phase C–M: not run.
Desktop	(Pentium	Q4_0 quant sweep (full, 5	Q4_K_M: repeatedly restarted,
G5420)		threads). Q3_K_M (11/12	never reached thread > 4. Phase
		rows).	B–M: not run.
Jetson Nano		Phase A, both target models,	Phase B (f16 and quantized KV):
		3-thread sweep (full).	0 rows, file present but empty.
Raspberry Pi 3B		Nothing.	Output file empty; the process
			did not complete a single bench-
			mark row.
Raspberry Pi 4B / 5		Not attempted.	Never part of the local hardware
			fleet.

Table 5.1: Data coverage by machine and phase; every unattributed claim in this thesis is drawn only from rows marked complete.

Concretely, this coverage implies the following specific boundaries on the current

model:

- **Hardware Fleet Scope.** Three devices produced usable data (Ryzen 5600H, Pentium G5420, Jetson Nano); a fourth (Raspberry Pi 3B) was attempted but produced none. This is not enough to fit a hardware-spec-to-parameters mapping with any statistical confidence – the per-device calibration constants in Section 4.1 are exactly that, per-device, not yet a general function of core count / bandwidth / ISA. Expanding to a broader, more standardized fleet (including a working Raspberry Pi run, a modern ARM board with SVE/SME, and an AVX-512 x86 chip) is the top priority for turning this into a genuine zero-shot predictor.
- **Serving and Concurrency.** Every measurement in this thesis is single-user, batch-size-1 generation. Phases F, G, I, and K1 – which targeted CPU pinning, memory-mapping, multi-process contention, and speculative decoding – did not produce usable data on any machine (Table 5.1), so the model says nothing yet about concurrent-request serving.
- **Quality Axis Calibration.** The perplexity deltas used to define the quality axis of the Pareto front are currently calibrated using established literature values for the GGUF ecosystem [21]. Completing local perplexity sweeps across all tested quantization formats will further ground this axis in hardware-specific measurements.
- **Contention and Scheduling.** The USL coefficients in Section 3.1 reflect default OS scheduling. Exploring pinned/isolated configurations and NUMA interleaving on multi-socket or chiplet-based systems would refine the coherency penalty (κ) for production deployment scenarios.
- **Flash Attention on Other Hardware.** Our Flash Attention findings in Section 3.3 are specific to the Ryzen 5600H. It is possible that on hardware with smaller L3 caches (e.g., the Pentium or Jetson), Flash Attention’s benefit could outweigh its overhead. This is an open question for future work.
- **KV-Cache Quantization on Bandwidth-Starved Hardware.** Our KV-cache quantization findings in Section 3.4 are specific to the compute-rich Ryzen 5600H. The roofline model predicts a sign-flip on bandwidth-starved hardware, but we have not yet tested this. This is a high-priority next experiment.

5.2 Future Work

The limitations above split naturally into two kinds: things this study already knows how to do and simply ran out of time for, and things that would require a materially larger or longer-running effort. This section addresses both, ordered from nearest-term to most ambitious.

5.2.1 Near-Term: Completing the Current Study

Building directly on the roofline model and Pareto framework, the following concrete next steps would close the specific gaps identified in Section 5.1:

1. **Expanding the hardware fleet** to include standardized edge and server platforms, enabling the regression of USL parameters directly against physical hardware specifications for zero-shot parameter prediction. Priority targets: a working Raspberry Pi 4/5 run, an x86 CPU with AVX-512 VNNI, an Apple Silicon Mac, and a multi-socket server with NUMA.
2. **Executing comprehensive local empirical benchmarks of diverse MoE architectures** to directly validate and refine the bytes-touched-per-token model, particularly focusing on router overhead and shared-layer bottlenecks. The staged models (Qwen1.5-MoE-A2.7B, DeepSeek-V2-Lite, DeepSeek-Coder-V2-Lite) will be benchmarked against dense equivalents on the same hardware.
3. **Completing local perplexity sweeps** to replace literature-based quality calibrations with hardware-specific measurements using `llama-perplexity` on the WikiText-2 dataset.
4. **Investigating the impact of advanced CPU scheduling**, CPU pinning, and NUMA interleaving on the coherency penalties observed during memory-bandwidth-bound generation.
5. **Testing Flash Attention and KV-cache quantization on bandwidth-starved hardware** (Pentium, Jetson) to validate the roofline model’s predicted sign-flips.

5.2.2 Longer-Term Directions

The near-term items above are things this study already knows how to do. This subsection takes a longer view: assuming they are completed and the roofline-USL model holds up, where could this line of work go next? The following seven directions are ordered roughly from nearest-term extensions of the existing framework to more ambitious departures from it.

From per-device calibration to a true zero-shot predictor

The central honest limitation of this thesis (Section 5.1) is that σ , κ , and BW_{eff} are calibrated separately for each device from a handful of local `llama-bench` runs, not derived from a device’s raw specifications. Turning this into a genuine zero-shot predictor – one that takes “6 cores, 2-channel DDR5, AVX-512” as input and outputs calibration constants without ever touching that hardware – requires calibrating enough devices that a regression against raw specs (core count, cache sizes, memory channels, ISA extensions) becomes statistically meaningful. Three partially-sampled devices, as this thesis has, is nowhere near enough; a follow-on study with on the order of twenty to thirty calibrated

devices, spanning several CPU vendors and generations, is the natural next scale for this work.

A living, community-contributed calibration dataset

Reaching that scale through one researcher’s personal hardware fleet is impractical, which is precisely the gap that led earlier drafts of this thesis to lean on external benchmark numbers that turned out to be difficult or impossible to verify (Section 3.8). A more sustainable path is to package the resumable calibration harness described in Section 2.6 as a lightweight, opt-in “auto-tune” companion to `llama.cpp` – something a user could run once after installation to get personalized flag recommendations for their own machine, with the option to contribute the resulting (anonymized) calibration constants back to a shared, versioned public dataset. This would directly address the verifiability problem this thesis ran into: every data point would come with a known provenance, a known harness version, and a known schema, rather than being a number copied from a blog post or forum thread.

From single-user batch-1 to concurrent serving

Every measurement in this thesis assumes a single user generating one response at a time (batch size 1), which is the right assumption for a personal chat application but not for a small on-premises server handling several users at once. Phases F and I of the instrumentation (Chapter 2) were designed to probe exactly this – CPU pinning, NUMA interleaving, and multi-process contention – but none produced usable data in this round (Table 5.1). Completing them is listed as near-term future work above; the longer-term direction is extending Equation 4.2’s roofline itself to a concurrent-request setting, where the bandwidth ceiling is shared across simultaneous generations rather than claimed by one, and the USL coherency term κ plausibly needs a second contention source (inter-process, not just inter-thread) added to it.

Generalizing the bytes-per-token lens beyond MoE

The MoE bandwidth argument in Chapter 4 rests on a single mechanism: a technique that reduces how many bytes must move through memory per generated token should, on memory-bound CPU hardware, produce a proportional throughput gain. Mixture-of-Experts routing is one way to reduce that byte count, but not the only one. Structured pruning, layer-skipping or early-exit techniques, and the small draft model used in speculative decoding (staged as Phase K1 in this study but never completed) all reduce effective bytes-touched-per-token through different mechanisms. A unified “effective active parameters” framework that treats MoE routing, pruning, and speculative decoding as three instances of the same underlying lever would let the roofline model predict the CPU speedup from any of them without re-deriving the argument from scratch each time.

Modeling conversations, not snapshots

Every context-depth result in this thesis (Sections 3.3 and 3.4) is a snapshot: throughput measured at one fixed context length. A real chat session is not a snapshot – it is a trajectory, with the KV cache growing turn by turn over what might be a fifty-turn conversation. Because both the Flash Attention regression and the KV-cache quantization sign-flip are context-length-dependent, a configuration that is optimal at turn one of a conversation is not guaranteed to still be optimal at turn thirty. Extending the Pareto front (Section 4.5) from a single-point prediction to a trajectory – predicted throughput as a function of turns elapsed, for a given configuration – would let a practitioner choose settings based on the conversation lengths they actually expect, rather than a single benchmark point that may not represent real usage.

Joint performance-energy rooflines for edge deployment

This thesis measured throughput exclusively; it did not measure power draw or thermal throttling, even though two of its four target devices (the Jetson Nano and the Raspberry Pi 3B) are exactly the class of hardware where power budget and sustained thermal output often matter as much as raw speed – a battery-powered or fanless deployment can be throughput-optimal on paper and still be unusable in practice if it draws too much power or throttles after ninety seconds. A natural extension of the roofline model already in this thesis is a second, joint roofline in performance-per-watt rather than raw performance, which would very likely reorder some of the recommendations in Section 4.6 – the fastest configuration and the most power-efficient one are not always the same configuration.

Testing how much of this is llama.cpp-specific

Every finding in this thesis – the Flash Attention regression, the Q4_0-versus-K-quant kernel effect, the USL coefficients themselves – was measured on one inference engine. Section 3.5 explicitly flags that we do not know whether the quantization-format anomaly reflects a property of the hardware and the algorithm, or an artifact of how llama.cpp specifically dispatches its CPU kernels. Repeating a subset of this study’s measurements on at least one other CPU inference engine (for example, ONNX Runtime or MLC-LLM) would separate genuine hardware/algorithm physics from implementation-specific quirks, and would tell us how much of the roofline-USL framework developed here is a statement about CPU inference in general versus a statement about this one, widely-used but not unique, engine.

Conclusion

This chapter has tried to state plainly what Chapters 3 and 4 do and do not support: a real, reproducible, mechanistically-explained account of thread scaling and memory-bandwidth effects on three specific CPUs, wrapped around a genuinely untested central

hypothesis about MoE models and calibrated on a fleet too small to generalize without further work. Neither half of that summary should be read as undermining the other – the measured effects stand on their own regardless of whether the MoE prediction holds up, and the model’s honesty about its own limits is what makes the near-term work above tractable rather than open-ended. The final chapter closes the thesis by drawing these threads together.

Chapter 6

Conclusion

Four consumer and embedded CPUs – three of which produced usable data – benchmarked with a resumable `llama.cpp` harness, were enough to recover four concrete and reproducible effects governing CPU-only LLM inference: (1) compute-bound prefill that saturates near the physical core count and follows Gunther’s Universal Scalability Law; (2) memory-bound generation that can retrograde as threads are added, most dramatically on bandwidth-starved processors; (3) quantization choices (both weight and KV-cache format) whose speed impact depends on which side of the compute/bandwidth boundary a given device and context length fall on; and (4) Flash Attention, which is a performance regression on CPU at long context. We formalized these effects into a roofline-style throughput model with per-device calibration constants, rather than a single universal equation, because the data honestly available does not yet support more than that. The model’s most interesting prediction – that Mixture-of-Experts architectures should be structurally advantaged on exactly the bandwidth-starved hardware this study focused on – remains directly untested on our own hardware and is loosely, not strongly, corroborated by outside sources; validating it is the clearest next step for turning this from a benchmarking exercise into the MoE-on-CPU research program it was originally motivated by.

Bibliography

- [1] Zhen Bi, Xueshu Chen, Luoyang Sun, Yuhang Yao, Qing Shen, Jungang Lou, and Cheng Deng. RooflineBench: A benchmarking framework for on-device LLMs via roofline analysis. *arXiv preprint arXiv:2602.11506*, 2026.
- [2] SMEPilot: Characterizing and optimizing LLM inference with scalable matrix extensions. *arXiv preprint arXiv:2606.16332*, 2026.
- [3] HeteroMosaic: Exposing and exploiting heterogeneous execution opportunities for energy-efficient edge LLM inference. *arXiv preprint arXiv:2607.12839*, 2026.
- [4] Which quantization should I use? a unified evaluation of llama.cpp quantization on Llama-3.1-8B-Instruct. *arXiv preprint arXiv:2601.14277*, 2026.
- [5] r/LocalLLaMA. Extensive LLama.cpp benchmark & more speed on CPU, 7b to 30b. <https://www.reddit.com/r/LocalLLaMA/comments/14ilo0t/>, 2023.
- [6] Woosuk Kwon et al. Efficient memory management for large language model serving with PagedAttention. In *SOSP*, 2023.
- [7] MLC-LLM: Universal LLM deployment engine. <https://github.com/mlc-ai/mlc-llm>.
- [8] Q4 vs Q5 vs Q6 vs Q8 quantization: Real quality loss numbers for local LLMs. <https://runaihome.com/blog/quantization-q4-q5-q6-q8-quality-loss-2026/>, 2026.
- [9] Samuel Williams, Andrew Waterman, and David Patterson. Roofline: An insightful visual performance model for multicore architectures. *Communications of the ACM*, 2009.
- [10] Neil J. Gunther. *Guerrilla Capacity Planning: A Tactical Approach to Planning for Highly Scalable Applications and Services*. Springer, 2007.
- [11] Elias Frantar et al. GPTQ: Accurate post-training quantization for generative pre-trained transformers. In *ICLR*, 2023.
- [12] Ji Lin et al. AWQ: Activation-aware weight quantization for LLM compression and acceleration. In *MLSys*, 2024.
- [13] ggml-org. llama.cpp perplexity tool README (quantization vs. perplexity reference table). <https://github.com/ggml-org/llama.cpp/blob/master/tools/perplexity/README.md>.

- [14] Mistral AI. Mixtral of experts, 2024.
- [15] DeepSeek-AI. DeepSeek-V2: A strong, economical, and efficient mixture-of-experts language model, 2024.
- [16] Qwen Team. Qwen1.5-MoE: A strong, economical, and efficient mixture-of-experts language model, 2024.
- [17] Jeff Geerling. LLMs accelerated with eGPU on a Raspberry Pi 5 / AI benchmarks. <https://github.com/geerlingguy/ai-benchmarks>, 2024.
- [18] LocalAIMaster. Raspberry Pi 5 LLM benchmarks: 12 models, real tokens. <https://localaimaster.com/blog/llm-raspberry-pi-5>, 2026.
- [19] GitHub / LocalLLaMA Community. KTransformers: DeepSeek-V3/R1 CPU inference benchmarks, 2024.
- [20] Tri Dao et al. FlashAttention: Fast and memory-efficient exact attention with IO-awareness. In *NeurIPS*, 2022.
- [21] Thomas Baruzier. Qwen2.5-7B-Instruct-GGUF: Perplexity and quantization metrics. <https://huggingface.co/ThomasBaruzier/Qwen2.5-7B-Instruct-GGUF>, 2024.